

About Dataset

This dataset contains information about used cars. This data can be used for a lot of purposes such as price prediction to exemplify the use of linear regression in Machine Learning.

The columns in the given dataset are as follows:

name

year

selling_price

km_driven

fuel

seller_type

transmission

Owner

<https://www.kaggle.com/datasets/nehalbirla/vehicle-dataset-from-cardekho>

Import libraries

```
In [87]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

Exploratory Data Analysis

```
In [88]: # Loading the car dataset
df = pd.read_csv("car data.csv")
df.head()
```

Out[88]:

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manual
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	Manual
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	Manual
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	Manual
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	Manual

In [89]: `# checking the no of rows and columns`
`df.shape`

Out[89]: (301, 9)

In [90]: `df.size`

Out[90]: 2709

In [91]: `# Getting info about dataset`
`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Car_Name        301 non-null    object
1   Year            301 non-null    int64
2   Selling_Price   301 non-null    float64
3   Present_Price   301 non-null    float64
4   Kms_Driven      301 non-null    int64
5   Fuel_Type       301 non-null    object
6   Seller_Type     301 non-null    object
7   Transmission    301 non-null    object
8   Owner           301 non-null    int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

In [92]: `df.describe().T`

Out[92]:

	count	mean	std	min	25%	50%	75%	max
Year	301.0	2013.627907	2.891554	2003.00	2012.0	2014.0	2016.0	2018.0
Selling_Price	301.0	4.661296	5.082812	0.10	0.9	3.6	6.0	35.0
Present_Price	301.0	7.628472	8.644115	0.32	1.2	6.4	9.9	92.6
Kms_Driven	301.0	36947.205980	38886.883882	500.00	15000.0	32000.0	48767.0	500000.0
Owner	301.0	0.043189	0.247915	0.00	0.0	0.0	0.0	3.0

In [93]: `# Checking the missing values`
`df.isna().sum()`

```
Out[93]: Car_Name      0
         Year        0
         Selling_Price  0
         Present_Price  0
         Kms_Driven    0
         Fuel_Type     0
         Seller_Type   0
         Transmission  0
         Owner        0
         dtype: int64
```

```
In [94]: # Checking the unique categories in each Categorical column
print(df['Fuel_Type'].unique())
print(df['Seller_Type'].unique())
print(df['Transmission'].unique())

['Petrol' 'Diesel' 'CNG']
['Dealer' 'Individual']
['Manual' 'Automatic']
```

```
In [95]: # Checking the distribution in each Categorical column
print(df.Fuel_Type.value_counts())

Fuel_Type
Petrol    239
Diesel    60
CNG        2
Name: count, dtype: int64
```

```
In [96]: print(df["Seller_Type"].value_counts())

Seller_Type
Dealer      195
Individual  106
Name: count, dtype: int64
```

```
In [97]: print(df["Transmission"].value_counts())

Transmission
Manual      261
Automatic   40
Name: count, dtype: int64
```

```
In [98]: # Checking no of duplicate in each column
df[df.duplicated()]
```

```
Out[98]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission
17	ertiga	2016	7.75	10.79	43000	Diesel	Dealer	Manual
93	fortuner	2015	23.00	30.61	40000	Diesel	Dealer	Automatic

```
In [99]: # Remove duplicates
df = df.drop_duplicates()
```

```
In [100]: df[df.duplicated()].sum()
```

```
Out[100]: Car_Name      0
          Year        0
          Selling_Price 0.0
          Present_Price 0.0
          Kms_Driven    0
          Fuel_Type     0
          Seller_Type    0
          Transmission  0
          Owner         0
          dtype: object
```

```
In [101... df.shape
```

```
Out[101]: (299, 9)
```

```
In [102... df.size
```

```
Out[102]: 2691
```

Encoding the categorical column

```
In [103... # Encoding "Fuel_Type" column
df.replace({'Fuel_Type' : {'Petrol':0, 'Diesel':1, 'CNG':2}}, inplace = True)

# Encoding "Seller_Type" column
df.replace({'Seller_Type' : {'Dealer':0, 'Individual':1}}, inplace = True)

# Encoding "Transmission" column
df.replace({'Transmission' : {'Manual':0, 'Automatic':1}}, inplace = True)
```

```
In [104... df.head()
```

```
Out[104]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission
0	ritz	2014	3.35	5.59	27000	0	0	0
1	sx4	2013	4.75	9.54	43000	1	0	0
2	ciaz	2017	7.25	9.85	6900	0	0	0
3	wagon r	2011	2.85	4.15	5200	0	0	0
4	swift	2014	4.60	6.87	42450	1	0	0

Train Test Split

```
In [105... # Split df into independent and Target variables
X = df.drop(columns=['Car_Name', 'Selling_Price'], axis=1)
Y = df['Selling_Price']
```

```
In [106... print(X)
```

	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	\
0	2014	5.59	27000	0	0	0	
1	2013	9.54	43000	1	0	0	
2	2017	9.85	6900	0	0	0	
3	2011	4.15	5200	0	0	0	
4	2014	6.87	42450	1	0	0	
..	
296	2016	11.60	33988	1	0	0	
297	2015	5.90	60000	0	0	0	
298	2009	11.00	87934	0	0	0	
299	2017	12.50	9000	1	0	0	
300	2016	5.90	5464	0	0	0	

	Owner
0	0
1	0
2	0
3	0
4	0
..	...
296	0
297	0
298	0
299	0
300	0

[299 rows x 7 columns]

In [107...

`print(Y)`

0	3.35
1	4.75
2	7.25
3	2.85
4	4.60
..	...
296	9.50
297	4.00
298	3.35
299	11.50
300	5.30

Name: Selling_Price, Length: 299, dtype: float64

In [113...

```
# import train_test_split from model_selection module
from sklearn.model_selection import train_test_split
```

In [211...

```
X_train,X_test,Y_train,Y_test = train_test_split(X, Y, test_size = 0.1, random_stat
```

Building a ML Model

1. Linear Regression Model

In [212...

```
# Import Linear Regression Model Library
from sklearn.linear_model import LinearRegression
```

In [213...

```
# Initilized Linear regression Model
lin_reg = LinearRegression()
```

In [214...

```
lin_reg.fit(X_train,Y_train)
```

Out[214]:

▼ LinearRegression

LinearRegression()

```
In [215... # Evaluate the model
from sklearn import metrics


# Prediction on training dataset
train_predict = lin_reg.predict(X_train)
```

```
In [216... # R square Error
R2_score_train = metrics.r2_score(Y_train, train_predict)
print("R squared Error:", R2_score_train)
```

R squared Error: 0.8771123681040541

Visualize the Actual and Predicted Prices(Train Set)

```
In [217... plt.scatter(Y_train, train_predict)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Prices vs Predicted Prices")
plt.show()
```




```
In [218... # Prediction on training dataset
test_predict = lin_reg.predict(X_test)
```

```
In [219... # R square Error
R2_score_test = metrics.r2_score(Y_test, test_predict)
print("R squared Error:", R2_score_test)
```

R squared Error: 0.8562516649198026

Visualize the Actual and Predicted Prices(Test Set)

```
In [220... plt.scatter(Y_test, test_predict)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Prices vs Predicted Prices")
plt.show()
```



2. Lasso Regression Model

```
In [221... # Import Lasso Regression Model Library
from sklearn.linear_model import Lasso
```

```
In [222... # Initilized Lasso regression Model
lasso_reg = Lasso()
```

```
In [223... lasso_reg.fit(X_train,Y_train)
```

```
Out[223]: ▾ Lasso
Lasso()
```


```
In [224... # Prediction on training dataset
train_predict = lasso_reg.predict(X_train)
```

```
In [225... # R square Error
R2_score_train = metrics.r2_score(Y_train, train_predict)
print("R squared Error:",R2_score_train)
```

R squared Error: 0.8428090658179708

Visualize the Actual and Predicted Prices(Train Set)

```
In [226... plt.scatter(Y_train, train_predict)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Prices vs Predicted Prices")
plt.show()
```




```
In [227... # Prediction on training dataset
test_predict = lasso_reg.predict(X_test)
```

```
In [228... # R square Error
R2_score_test = metrics.r2_score(Y_test, test_predict)
print("R squared Error:", R2_score_test)
```

R squared Error: 0.7933398240755692

Visualize the Actual and Predicted Prices(Test Set)

```
In [229... plt.scatter(Y_test, test_predict)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Prices vs Predicted Prices")
plt.show()
```

In []:

In []: